

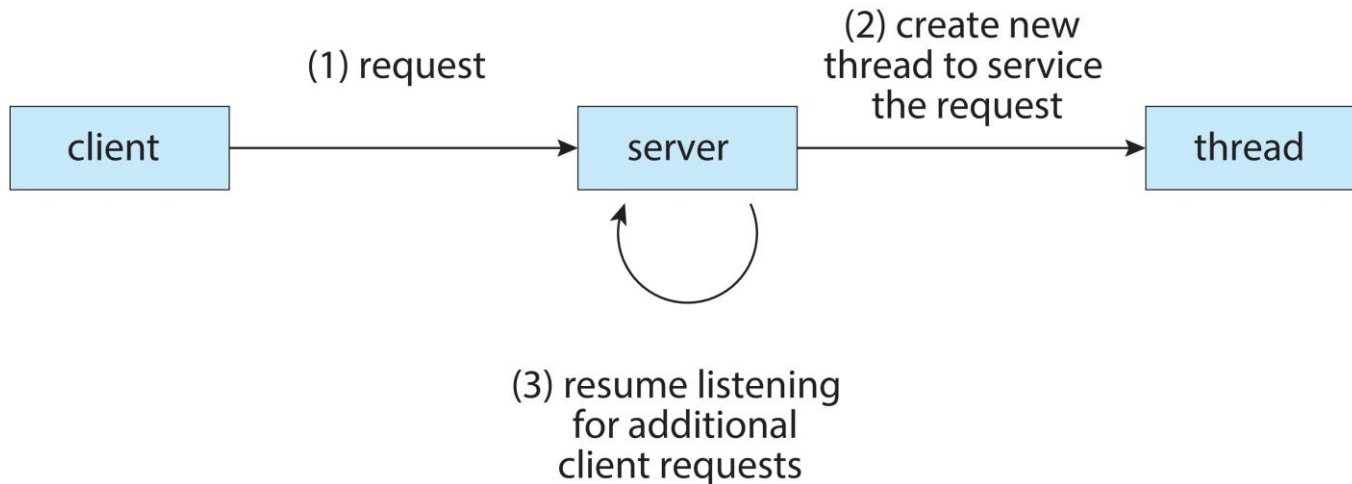
Chapter 4: Threads & Concurrency





Motivation

- Most modern applications are multithreaded
 - E.g., web browser, word processor, web server
- Multiple tasks within the application can be implemented by separate threads
 - E.g., update display, fetch data, spell checking, answer a network request



Multithreaded server architecture





What is a Thread?

- A **thread** is a basic unit of CPU utilization
- A process can contain multiple threads of control
- A thread consists of
 - A thread ID
 - A program counter
 - A register set
 - A stack
- Threads belonging to the same process share
 - code section, data section, heap
 - open files

Linux command to display threads in a process


```
$ top -H -p <pid>
```

```
$ ps -T -p <pid>
```






Single and Multithreaded Processes



single-threaded process



multithreaded process





Thread Control Block (TCB)

- OS maintains a TCB for each thread
- TCB contains
 - Thread ID
 - Stack pointer
 - Program counter
 - Register values
 - Pointer to PCB of the process that the thread belongs to





Benefits of Multithreading

- **Responsiveness** – may allow continued execution of a program if part of it is blocked or is performing a lengthy operation; especially important for user interfaces
- **Resource Sharing** – threads share resources of process, sharing data easier than shared memory or message passing
- **Economy** – thread creation is cheaper (i.e., consumes less time & memory) than process creation, thread context-switching has lower overhead than process-context switching
 - Thread context-switching involves saving/loading registers and program counter to/from TCB; no need to save/load memory management information
- **Scalability** – multithreaded process can take advantage of multicore architectures
 - Threads may be running in parallel on different processing cores





Multicore Programming

- Recent trend in computer design is to place multiple computing cores on a single processing chip – **multicore** systems
- Multithreaded programming provides more efficient use of multiple computing cores and improved concurrency
- **Parallelism** implies a system can perform more than one task simultaneously
- **Concurrency** supports more than one task making progress
 - Single processor/core: scheduler providing concurrency
 - It is possible to have concurrency without parallelism







Concurrency vs. Parallelism

Consider an application with 4 threads:

□ Concurrent execution on single-core system:



□ Parallelism on a multi-core system:





Amdahl's Law

- Identifies performance gains from adding additional computing cores to an application that has both serial and parallel components
- S is serial portion of application, N processing cores

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$


- E.g., if application is 75% parallel & 25% serial, moving from 1 to 3 cores results in speedup of 2 times
- As N approaches infinity, speedup approaches $1/S$

Serial portion of an application has disproportionate effect on performance gained by adding additional cores





Amdahl's Law





User Threads and Kernel Threads

- ❑ **User threads** – supported at the user level, management done by user-level threads library
- ❑ **Kernel threads** – supported by the kernel
- ❑ Virtually all contemporary operating systems support kernel threads:
 - ❑ Windows
 - ❑ Linux
 - ❑ macOS
 - ❑ iOS
 - ❑ Android


Linux command to see Kernel threads
\$ ps -ef





Multithreading Models


- Relationship between user threads and kernel threads
 - Many-to-One
 - One-to-One
 - Many-to-Many





Many-to-One


- ❑ Many user threads mapped to one kernel thread
- ❑ Thread management done by thread library in user space → efficient
- ❑ Entire process will block if a thread makes a blocking system call
- ❑ Multiple threads cannot run in parallel on multicore system
- ❑ Few systems currently use this model





One-to-One


- ❑ Each user thread is mapped to a kernel thread
- ❑ Greater concurrency than many-to-one: another thread can run when a thread makes a blocking system call
- ❑ Multiple threads can run in parallel on multicore system
- ❑ Creating a user thread requires creating a kernel thread → expensive
 - ❑ Number of threads per process sometimes restricted due to overhead
- ❑ Most operating systems now use one-to-one model





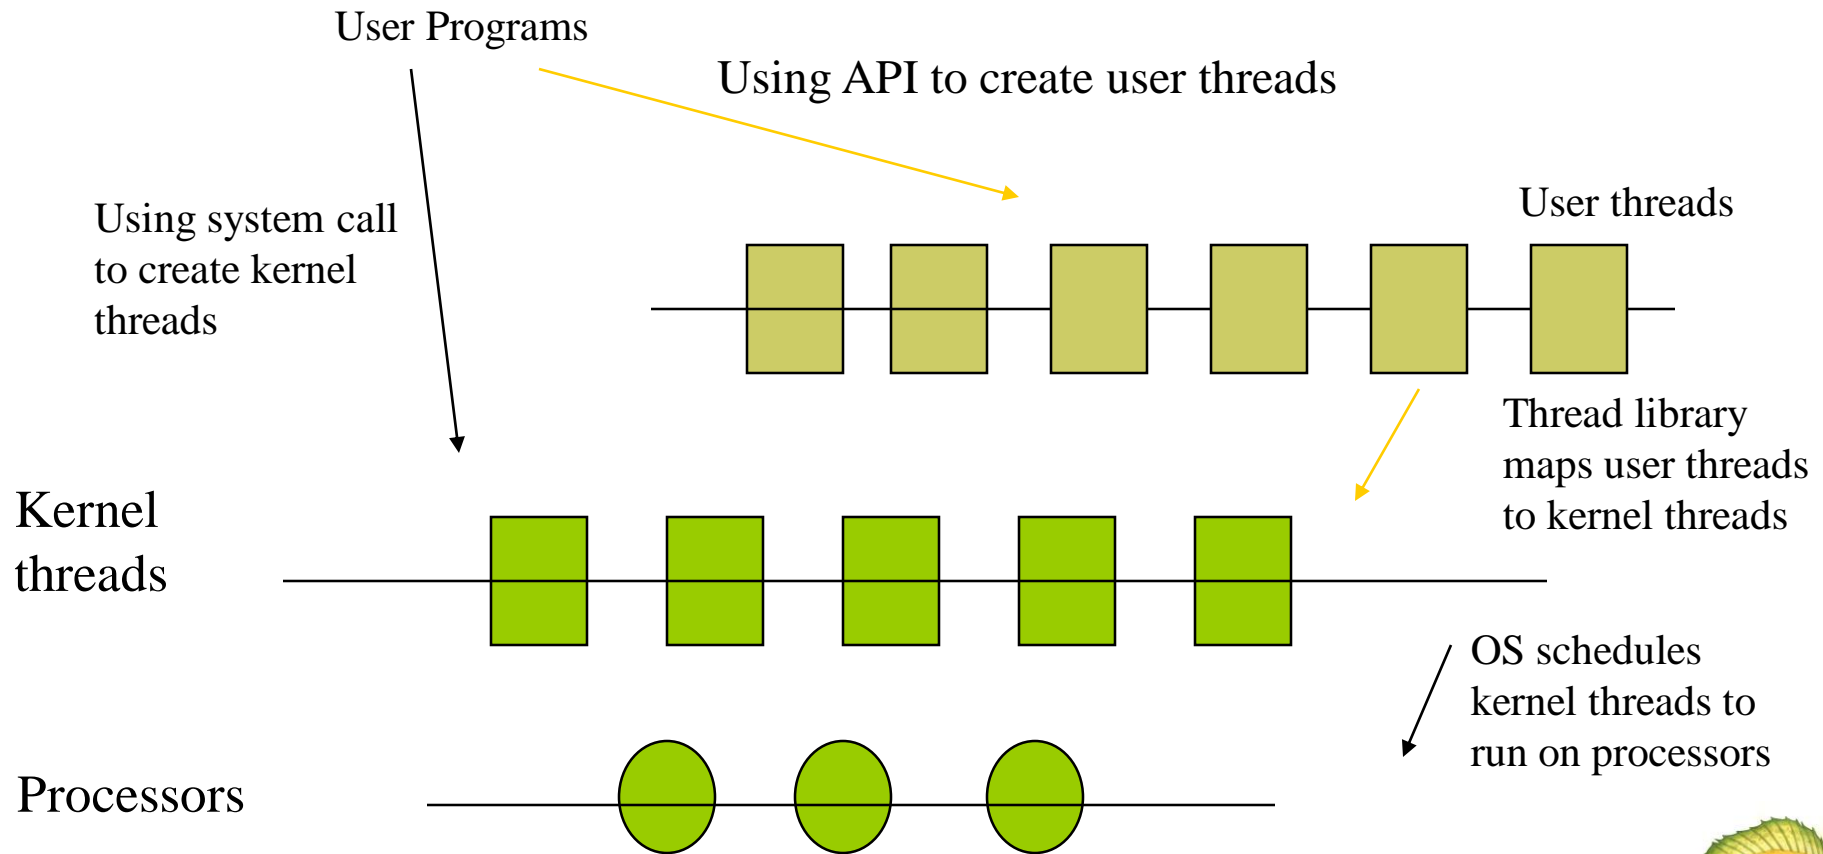
Many-to-Many Model

- ❑ Many user-level threads are multiplexed to a smaller or equal number of kernel threads
- ❑ Application developer can create as many user threads as necessary, corresponding kernel threads can run in parallel on a multicore system
- ❑ When a user thread performs a blocking system call, the kernel can schedule another thread for execution
- ❑ Not very common





Threads in a Computer System





Thread Libraries

- **Thread library** provides programmer with API for creating and managing threads
- Two primary ways of implementing thread library
 - Library entirely in user space: invoking a library function results in a local function call in user space
 - Kernel-level library supported by the OS: invoking a library function results in a system call to kernel
- Three main thread libraries:
 - POSIX pthreads: can be user level or kernel level
 - Windows threads: kernel level
 - Java threads: implemented using a thread library available on the host system





Pthreads

- pthreads is a Portable Operating System Interface (POSIX) standard API for thread creation and synchronization
 - API specifies behavior of the thread library, implementation is up to development of the library
 - Available on many UNIX-like operating systems (Solaris, Linux, macOS)

- Using pthreads on pyrite:
 - #include <pthread.h>
 - Compile and link with the pthread library: `gcc pthreads.c -lpthread`





pthread_create (1)

int pthread_create(pthread_t *thread, pthread_attr_t *attr,
void * (* start_routine) (void *), void *arg) – create a new thread

- thread points to a buffer that stores the ID of the new thread
- attr points to a structure containing the attributes of the new thread
 - If attr is NULL, the thread is created with default attributes
- The new thread starts execution by invoking start_routine()
- arg is a pointer to the argument of start_routine()
 - If multiple arguments are needed, arg points to a data structure that contains all arguments
- Returns 0 on success, returns an error number on error





pthread_create (2)

- ❑ The new thread executes concurrently with the parent thread
- ❑ The new thread runs until one of the following happens
 - ❑ It returns from start routine
 - ❑ It calls **pthread_exit()**
 - ❑ Any of the threads in the process calls **exit()** or the main thread performs a return from **main()**. This causes the termination of all threads in the process.





pthread_create() Example

```
#include <pthread.h>
```

```
#include <stdio.h>
```

```
#include <unistd.h>
```

```
void *my_thread (void *arg)
{
    char *msg = (char *) arg;
    printf("Thread says %s\n", msg);
}
```

```
int main (int argc, char *argv[])
{
    pthread_t t;
    char msg[20] = "Hello World";
    pthread_create(&t, NULL, my_thread, msg);
    sleep(3); //what happens if this statement is removed?
    return 0;
}
```

<https://onlinegdb.com/HyTorJxED>





pthread_exit()

#include <pthread.h>

void pthread_exit(void *retval) – terminate calling thread

- The function returns a value via retval that is available to another thread in the same process that calls pthread_join()
- The function does not return to the caller





pthread_join()

#include <pthread.h>

int pthread_join(pthread_t th, void **retval) - wait for a thread to terminate

- ❑ th: the thread to wait for
- ❑ If retval is not NULL, then **pthread_join()** copies the exit status of the target thread into the location pointed to by retval
- ❑ Returns 0 on success, returns an error number on error
- ❑ When a thread terminates, its TCB is not deallocated until another thread performs **pthread_join()** on it





Pthread_join() Example

```
#include <pthread.h>
#include <stdio.h>
#include <unistd.h>

int N;

void * thread(void *x)
{
    int *id = (int *)x;
    while (N != *id);
    printf("Thread %d \n", *id);
    N--;
    pthread_exit(NULL);
}

int main()
{
    N = 0;
    int id1=1;
    int id2=2;
    int id3=3;

    pthread_t t1, t2, t3;
    printf("Parent creating threads\n");
    pthread_create(&t1, NULL, thread, &id1);
    pthread_create(&t2, NULL, thread, &id2);
    pthread_create(&t3, NULL, thread, &id3);
    printf("Threads created\n");
    N = 3;
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    pthread_join(t3, NULL);
    printf("Threads are done\n");
    return 0;
}
```

