


# Chapter 8: Deadlocks

Chapter 8: Section 8.1-8.6





# System Model

- System consists of resources
- Resource types  $R_1, R_2, \dots, R_m$ 
  - Physical resource types: CPUs, memory, I/O devices
  - Logical resource types: files, mutex locks, semaphores
- Each resource type  $R_i$  has one or more identical instances
- Each thread utilizes a resource as follows:
  - **Request** – If the request cannot be granted immediately, then the thread must wait until it can acquire the resource
  - **Use**
  - **Release**





# What is a Deadlock?

- A **deadlock** is a situation involving a set of threads in which each thread waits for an event that can be caused only by another thread in the set
- Deadlock example when using semaphores

Two threads share semaphores A and B, both initialized to 0

$T_1$

wait (A);

...


signal(B);

$T_2$

wait(B);

...

signal(A);





# Deadlock when Using Mutex Locks

- Two mutex locks are created and initialized:

```
pthread_mutex_t first_mutex;  
pthread_mutex_t second_mutex;  
  
pthread_mutex_init(&first_mutex, NULL);  
pthread_mutex_init(&second_mutex, NULL);
```

- Deadlock is possible if thread 1 acquires **first\_mutex** and thread 2 acquires **second\_mutex**. Thread 1 then waits for **second\_mutex** and thread 2 waits for **first\_mutex**.

```
/* thread_one runs in this function */  
void *do_work_one(void *param)  
{  
    pthread_mutex_lock(&first_mutex);  
    pthread_mutex_lock(&second_mutex);  
    /**  
     * Do some work  
     */  
    pthread_mutex_unlock(&second_mutex);  
    pthread_mutex_unlock(&first_mutex);  
  
    pthread_exit(0);  
}
```

```
/* thread_two runs in this function */  
void *do_work_two(void *param)  
{  
    pthread_mutex_lock(&second_mutex);  
    pthread_mutex_lock(&first_mutex);  
    /**  
     * Do some work  
     */  
    pthread_mutex_unlock(&first_mutex);  
    pthread_mutex_unlock(&second_mutex);  
  
    pthread_exit(0);  
}
```





# Deadlock Characterization

Deadlock occurs when four conditions hold simultaneously

- **Mutual exclusion:** At least one resource must be held in a non-sharable mode; that is, only one thread at a time can use the resource
- **Hold and wait:** A thread must be holding at least one resource and waiting to acquire additional resources that are currently being held by other threads
- **No preemption:** A resource can be released only voluntarily by the thread holding it, after that thread has completed its task
- **Circular wait:** There exists a set  $\{T_0, T_1, \dots, T_n\}$  of waiting threads such that  $T_0$  is waiting for a resource that is held by  $T_1$ ,  $T_1$  is waiting for a resource that is held by  $T_2$ , ...,  $T_{n-1}$  is waiting for a resource that is held by  $T_n$ , and  $T_n$  is waiting for a resource that is held by  $T_0$ .






# Resource-Allocation Graph

- We can use resource-allocation graphs to study deadlocks
- A resource-allocation graph has a set of vertices  $V$  and a set of edges  $E$
- $V$  is partitioned into two types:

- $T = \{T_1, T_2, \dots, T_n\}$ , the set consisting of all threads in the system




- $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system





- $E$  is partitioned into two types:

- **Request edge** – directed edge  $T_i \rightarrow R_j$  indicating thread  $T_i$  has requested an instance of resource type  $R_j$




- **Assignment edge** – directed edge  $R_j \rightarrow T_i$  indicating an instance of resource type  $R_j$  has been allocated to thread  $T_i$







# Resource Allocation Graph Example

- 1 instance of R1
- 2 instances of R2
- 1 instance of R3
- 3 instances of R4
- T1 holds one instance of R2 and is waiting for an instance of R1
- T2 holds one instance of R1, one instance of R2, and is waiting for an instance of R3
- T3 holds one instance of R3





# Resource Allocation Graph with No Deadlock




If graph contains no cycles, then there is no deadlock







# Resource Allocation Graph With A Deadlock




- Processes  $P_1$ ,  $P_2$ , and  $P_3$  are deadlocked
- Does a cycle imply a deadlock?  
Answer is No!





# Graph With A Cycle But No Deadlock



When  $P_4$  releases an instance of  $R_2$ , that resource can be allocated to  $P_3$ , breaking the cycle





# Basic Facts

---

- If a resource-allocation graph contains no cycles  $\Rightarrow$  no deadlock
- If a resource-allocation graph contains a cycle, then the system **may** or **may not** be in a deadlocked state
  - If there is only one instance per resource type  $\Rightarrow$  deadlock
  - If there are several instances per resource type  $\Rightarrow$  possibility of deadlock





# Methods for Handling Deadlocks

---

- Ensure that the system will **never** enter a deadlock state:
  - Deadlock prevention
  - Deadlock avoidance
- Deadlock detection and recovery:
  - Allow the system to enter a deadlock state
  - Abort a process or preempt some resources when a deadlock is detected
- Ignore the problem all together
  - Simply let deadlocks happen and reboot as necessary
  - Used by most operating systems, including Linux and Windows





# Deadlock Prevention

---

Ensure that one of the four necessary conditions for deadlock cannot hold

- **Mutual Exclusion** – not required for sharable resources (e.g., read-only files); must hold for non-sharable resources (e.g., mutex locks)
- **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
  - Require process to request and be allocated all its resources before it begins execution
  - Or allow process to request resources only when the process has none allocated to it
  - Drawbacks: low resource utilization, starvation possible





# Deadlock Prevention (Cont.)

## □ **No Preemption** –

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are preempted
- Preempted resources are added to the list of resources for which the process is waiting
- Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting
- Some resources can be preempted (e.g., CPU, memory space); some resources cannot be preempted (e.g., mutex locks)

## □ **Circular Wait** – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration





# Denying Circular Wait

- Simply assign each resource a unique number
- Resources must be acquired in an increasing order of enumeration
- For example:

**F(first\_mutex) = 1**

**F(second\_mutex) = 5**

code for **thread\_two** cannot  
be written as this: →

```
/* thread_one runs in this function */
void *do_work_one(void *param)
{
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);

    pthread_exit(0);
}

/* thread_two runs in this function */
void *do_work_two(void *param)
{
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);

    pthread_exit(0);
}
```





# Deadlock Avoidance

---

Requires that the system has some additional ***a priori*** information available

- Each process declares the ***maximum number*** of resources of each type that it may need
- Whenever a process requests a resource that is currently available, the system decides whether the resource can be allocated immediately or whether the process must wait
- The system makes the decision by considering the resource-allocation state of the system
  - **Resource-allocation state** is defined by the number of available and allocated resources, and the maximum demands of the processes







# Safe State


- When a process requests an available resource, system must decide if immediate allocation leaves the system in a **safe state**
- System is in **safe state** if there exists a safe sequence
  - A sequence of processes  $\langle P_1, P_2, \dots, P_n \rangle$  is a **safe sequence** for the current allocation state if, for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all  $P_j$ , with  $j < i$
- That is:
  - If  $P_i$ 's resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished
  - When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate
  - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on





# Basic Facts

- If a system is in safe state  $\Rightarrow$  no deadlocks
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock
- Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state





# Safe State Example

- Consider a system with 12 resources and 3 processes
- System state at  $t_0$ :

	<u>Maximum Needs</u>	<u>Current Allocation</u>	Available = 3
$P_0$	10	5	
$P_1$	4	2	
$P_2$	9	2	

The system is in a safe state at  $t_0$  because  $\langle P_1, P_0, P_2 \rangle$  is a safe sequence

- Suppose at  $t_1$   $P_2$  requests and is allocated one more resource

	<u>Maximum Needs</u>	<u>Current Allocation</u>	Available = 2
$P_0$	10	5	
$P_1$	4	2	
$P_2$	9	3	

The system is no longer in a safe state at  $t_1$





# Deadlock Avoidance Algorithms

---


- Single instance of a resource type
  - Use a resource-allocation graph
  
- Multiple instances of a resource type
  - Use the Banker's Algorithm





# Resource-Allocation Graph Scheme


- **Claim edge**  $T_i \rightarrow R_j$  indicates that thread  $T_i$  may request resource  $R_j$  at some time in the future
  - Claim edge represented by a *dashed line*
- Claim edge is converted to request edge when a thread requests a resource
- Request edge is converted to an assignment edge when the resource is allocated to the thread
- When a resource is released by a thread, assignment edge reconverts to a claim edge
- Resources must be claimed *a priori* in the system





# Resource-Allocation Graph Algorithm

- Suppose that thread  $T_i$  requests a resource  $R_j$
- The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource-allocation graph
- Example: Suppose  $T_2$  requests  $R_2$ 
  - Allocating  $R_2$  to  $T_2$  will create a cycle in the graph. Thus, the request cannot be granted.





# Banker's Algorithm

---

- Applicable to a system with multiple instances of each resource type
- Each process must a priori claim maximum use of each resource type
- When a process requests a resource that is available, it may have to wait





# Data Structures for the Banker's Algorithm

Let  $n$  = number of processes,  $m$  = number of resource types.

□ **Available:** Vector of length  $m$

□ If available  $[j] = k$ , there are  $k$  instances of resource type  $R_j$  available

Example: Available

A	B	C
---	---	---

3	5	4
---	---	---

□ **Max:**  $n \times m$  matrix

□ If  $Max[i, j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$

Example: Max

	A	B	C
--	---	---	---

$P_0$	3	4	0
-------	---	---	---

$P_1$	2	0	3
-------	---	---	---







# Data Structures for the Banker's Algorithm

□ **Allocation:**  $n \times m$  matrix.

□ If  $\text{Allocation}[i,j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$

Example:      Allocation

	A	B	C
$P_0$	2	3	0
$P_1$	2	0	1

□ **Need:**  $n \times m$  matrix.

□ If  $\text{Need}[i,j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task

$$\text{Need}[i,j] = \text{Max}[i,j] - \text{Allocation}[i,j]$$

Example:  $\text{Need}_0 = \text{Max}_0 - \text{Allocation}_0 = (3,4,0) - (2,3,0) = (1,1,0)$





# Safety Algorithm

The algorithm determines whether or not a system is in a safe state.

1. Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , respectively.  
Initialize:

**Work = Available**

**Finish** [ $i$ ] = **false** for  $i = 0, 1, \dots, n-1$

2. Find an  $i$  such that both:

(a) **Finish** [ $i$ ] = **false**

(b) **Need** <sub>$i$</sub>  ≤ **Work**

If no such  $i$  exists, go to step 4

3. **Work = Work + Allocation** <sub>$i$</sub>   
**Finish** [ $i$ ] = **true**  
go to step 2

4. If **Finish** [ $i$ ] == **true** for all  $i$ , then the system is in a safe state





# Example of Safety Algorithm (1)

- 5 processes  $P_0$  through  $P_4$ ;

3 resource types:

$A$  (10 instances),  $B$  (5 instances), and  $C$  (7 instances)

- Snapshot at time  $t_0$ :

	<u>Allocation</u>	<u>Max</u>	<u>Need</u>	<u>Available</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
$P_0$	0 1 0	7 5 3	7 4 3	3 3 2
$P_1$	2 0 0	3 2 2	1 2 2	
$P_2$	3 0 2	9 0 2	6 0 0	
$P_3$	2 1 1	2 2 2	0 1 1	
$P_4$	0 0 2	4 3 3	4 3 1	

- The system is in a safe state since there exists a safe sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$





# Example of Safety Algorithm (2)

	<u>Allocation</u>	<u>Max</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C	A B C
$P_0$	0 1 0	7 5 3	7 4 3	3 3 2
$P_1$	2 0 0	3 2 2	1 2 2	
$P_2$	3 0 2	9 0 2	6 0 0	
$P_3$	2 1 1	2 2 2	0 1 1	
$P_4$	0 0 2	4 3 3	4 3 1	

$\langle P_1, P_3, P_4, P_2, P_0 \rangle$  is a safe sequence as shown below:

Work = Available = (3, 3, 2)

1.  $P_1$  can finish because  $\text{Need}_1 = (1, 2, 2) \leq \text{Available} = (3, 3, 2)$ . When  $P_1$  finishes,  $\text{Available} = (3, 3, 2) + (2, 0, 0) = (5, 3, 2)$
2.  $P_3$  can finish because  $\text{Need}_3 = (0, 1, 1) \leq \text{Available} = (5, 3, 2)$ . When  $P_3$  finishes,  $\text{Available} = (5, 3, 2) + (2, 1, 1) = (7, 4, 3)$
3.  $P_4$  can finish because  $\text{Need}_4 = (4, 3, 1) \leq \text{Available} = (7, 4, 3)$ . When  $P_4$  finishes,  $\text{Available} = (7, 4, 3) + (0, 0, 2) = (7, 4, 5)$
4.  $P_2$  can finish because  $\text{Need}_2 = (6, 0, 0) \leq \text{Available} = (7, 4, 5)$ . When  $P_2$  finishes,  $\text{Available} = (7, 4, 5) + (3, 0, 2) = (10, 4, 7)$
5.  $P_0$  can finish because  $\text{Need}_0 = (7, 4, 3) \leq \text{Available} = (10, 4, 7)$ . When  $P_0$  finishes,  $\text{Available} = (10, 4, 7) + (0, 1, 0) = (10, 5, 7)$





# Resource-Request Algorithm

- Let  **$Request_i$**  = request vector for process  **$P_i$** . If  **$Request_i[j] = k$**  then process  **$P_i$**  wants  **$k$**  instances of resource type  **$R_j$** 
  - e.g.,  **$Request_1 = (1,0,2)$**
- When a request is made by  **$P_i$** , the following actions are taken to determine whether the request can be safely granted
  1. If  **$Request_i \leq Need_i$** , go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
  2. If  **$Request_i \leq Available$** , go to step 3. Otherwise  **$P_i$**  must wait, since resources are not available
  3. Pretend to allocate requested resources to  **$P_i$**  by modifying the state as follows:

$$Available = Available - Request_i;$$

$$Allocation_i = Allocation_i + Request_i;$$

$$Need_i = Need_i - Request_i;$$

- If new state is safe  $\Rightarrow$  the resources are allocated to  **$P_i$**
- If new state is unsafe  $\Rightarrow$   **$P_i$**  must wait, and the old resource-allocation state is restored





# Resource-Request Algorithm Example

- Suppose the current state of the system is the following:

	<u>Allocation</u>	<u>Max</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C	A B C
$P_0$	0 1 0	7 5 3	7 4 3	3 3 2
$P_1$	2 0 0	3 2 2	1 2 2	
$P_2$	3 0 2	9 0 2	6 0 0	
$P_3$	2 1 1	2 2 2	0 1 1	
$P_4$	0 0 2	4 3 3	4 3 1	

- We have shown that the system is in a safe state
- Suppose now that process  $P_1$  makes a request **Request<sub>1</sub>** = (1,0,2)





## Example: $P_1$ Requests (1,0,2)

- Check that  $\text{Request}_1 \leq \text{Need}_1$ , that is,  $(1,0,2) \leq (1,2,2) \Rightarrow \text{true}$
- Check that  $\text{Request}_1 \leq \text{Available}$ , that is,  $(1,0,2) \leq (3,3,2) \Rightarrow \text{true}$
- Pretend that the request is granted, then we arrive at the new state:
  - $\text{Available} = \text{Available} - \text{Request} = (3,3,2) - (1,0,2) = (2,3,0)$
  - $\text{Allocation}_1 = \text{Allocation}_1 + \text{Request} = (2,0,0) + (1,0,2) = (3,0,2)$
  - $\text{Need}_1 = \text{Need}_1 - \text{Request} = (1,2,2) - (1,0,2) = (0,2,0)$

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	2 3 0
$P_1$	3 0 2	0 2 0	
$P_2$	3 0 2	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	





# Example: $P_1$ Requests (1,0,2)

- The new system state is the following:

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	2 3 0
$P_1$	3 0 2	0 2 0	
$P_2$	3 0 2	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

- The new state is safe as we can find a safety sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  as shown below. Thus, we can grant the request.

Work = Available = (2,3,0)

1.  $P_1$  can finish because  $\text{Need}_1 = (0,2,0) \leq \text{Available} = (2,3,0)$ . When  $P_1$  finishes,  $\text{Available} = (2,3,0) + (3,0,2) = (5,3,2)$
2.  $P_3$  can finish because  $\text{Need}_3 = (0,1,1) \leq \text{Available} = (5,3,2)$ . When  $P_3$  finishes,  $\text{Available} = (5,3,2) + (2,1,1) = (7,4,3)$
3.  $P_4$  can finish because  $\text{Need}_4 = (4,3,1) \leq \text{Available} = (7,4,3)$ . When  $P_4$  finishes,  $\text{Available} = (7,4,3) + (0,0,2) = (7,4,5)$
4.  $P_0$  can finish because  $\text{Need}_0 = (7,4,3) \leq \text{Available} = (7,4,5)$ . When  $P_0$  finishes,  $\text{Available} = (7,4,5) + (0,1,0) = (7,5,5)$
5.  $P_2$  can finish because  $\text{Need}_2 = (6,0,0) \leq \text{Available} = (7,5,5)$ . When  $P_2$  finishes,  $\text{Available} = (7,5,5) + (3,0,2) = (10,5,7)$







## Example: $P_4$ Requests (3,3,0)

- Suppose process  $P_4$  makes a request  $\mathbf{Request}_4 = (3,3,0)$  when the system is in the following state. Can the request be granted?

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	2 3 0
$P_1$	3 0 2	0 2 0	
$P_2$	3 0 2	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

- Check that  $\mathbf{Request}_4 \leq \mathbf{Need}_4$ , that is,  $(3,3,0) \leq (4,3,1) \Rightarrow \text{true}$
- Check that  $\mathbf{Request}_4 \leq \mathbf{Available}$ , that is,  $(3,3,0) \leq (2,3,0) \Rightarrow \text{false}$
- Thus, the request cannot be granted as there are not enough resources to satisfy the request





## Example: $P_0$ Requests (0,2,0)

- Suppose process  $P_0$  makes a request **Request**<sub>0</sub> = (0,2,0) when the system is in the following state. Can the request be granted?

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	2 3 0
$P_1$	3 0 2	0 2 0	
$P_2$	3 0 2	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

- Check that  $\text{Request}_0 \leq \text{Need}_0$ , that is,  $(0,2,0) \leq (7,4,3) \Rightarrow \text{true}$
- Check that  $\text{Request}_0 \leq \text{Available}$ , that is,  $(0,2,0) \leq (2,3,0) \Rightarrow \text{true}$
- Next, we pretend that the request is granted, then we arrive at the new state:
  - $\text{Available} = \text{Available} - \text{Request} = (2,3,0) - (0,2,0) = (2,1,0)$
  - $\text{Allocation}_0 = \text{Allocation}_0 + \text{Request} = (0,1,0) + (0,2,0) = (0,3,0)$
  - $\text{Need}_0 = \text{Need}_0 - \text{Request} = (7,4,3) - (0,2,0) = (7,2,3)$





## Example: $P_0$ Requests (0,2,0)

- The new state is the following:

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 3 0	7 2 3	2 1 0
$P_1$	3 0 2	0 2 0	
$P_2$	3 0 2	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

- Now let's try to find a safe sequence:

Work = Available = (2,1,0)

We cannot find a process  $P_i$  such that  $\text{Need}_i \leq \text{Available}$ . Thus, no safe sequence exists.

Therefore, the system state is unsafe and the request cannot be granted!

